
Statistical Analysis of Retail Electricity Sales in the United States

A Multiregional, Multisector Time Series Evaluation (2014–2024)

Mohankrishna Kolla

December 5, 2025

Abstract

Abstract. This report investigates temporal, seasonal, sectoral, and regional variation in retail electricity sales across the United States using archival monthly energy data from 2014–2024. Using exploratory data analysis, grouped descriptive statistics, and multiple linear regression on log-transformed sales, the study evaluates structural differences across regions and sectors. Results indicate strong statistical significance in regional and sectoral variation ($R^2 \approx 0.852$), with the South region and Residential sectors driving peak demand. The Transportation sector exhibits minimal grid impact relative to other sectors. Implications for demand forecasting and energy planning are discussed.

Contents

1	Research Question and Motivation	3
2	Methods	3
2.1	Data Source and Preparation	3
3	Exploratory Data Analysis	4
3.1	Overall Summary Statistics	4
3.2	Overall Temporal Pattern	4
3.3	Distribution by Region and Sector	5
3.4	Normality Assessment	5
3.5	Regional Time Series	6
3.6	Sector-Specific Normality	7
4	Grouped Summary Statistics	8
5	Statistical Analysis	9
5.1	Regression Model Results	9
5.2	Interpretation	10
6	Discussion	11
6.1	The “South Effect” and Regional Heterogeneity	11
6.2	The Transportation Gap: A Looming Grid Shock	11
6.3	Economic Decoupling: The Absence of a Time Trend	12
6.4	Methodological Insights: The Power of Log-Normality	12
6.5	Strategic Recommendations	13
7	Conclusion	13
A	Python Source Code	15

1 Research Question and Motivation

The central research questions of this study are:

1. How do retail electricity sales differ across economic sectors (Residential, Commercial, Industrial, Transportation)?
2. How do these sectoral differences vary across U.S. Census regions (Northeast, Midwest, South, West)?
3. Is there evidence of a long-term national trend in electricity consumption over the 2014–2024 period?
4. How do seasonal patterns manifest in aggregate electricity demand?

Understanding these relationships is essential for long-term demand forecasting, infrastructure planning, and regional energy policy. Variation in electricity use reflects climate, demographics, regional industry, and transportation electrification.

2 Methods

2.1 Data Source and Preparation

This project uses publicly available archival monthly electricity sales data obtained from the U.S. Energy Information Administration (EIA). The dataset includes monthly observations of retail electricity sales (MWh) categorized by State and Sector.

Key preprocessing steps included:

- **Filtering:** Removal of aggregate "Total" rows to prevent double counting.
- **Feature Engineering:** Mapping states to Census Regions (Northeast, Midwest, South, West) and extracting seasonal indicators.
- **Transformation:** Application of a log-transformation ($\log(1 + \text{Value})$) to address heteroscedasticity and the heavy right-skew of electricity sales data.
- **Sampling:** A subset of 5,000 observations was used for the regression model to ensure computational efficiency while maintaining statistical power.

3 Exploratory Data Analysis

3.1 Overall Summary Statistics

Table 1 presents the summary statistics for the filtered dataset. The data exhibits high variance, with a standard deviation exceeding the mean in the original scale, necessitating the log transformation used in the inferential modeling.

Table 1: Summary statistics for retail electricity sales (Original vs. Log-Transformed).

Variable	Count	Mean	Std. Dev.	Min	Max
Original Value (MWh)	25,695	1,325.93	1,489.76	0	8,499.00
Log-Transformed	25,695	5.54	2.85	0	9.05

3.2 Overall Temporal Pattern

Figure 1 shows the mean electricity sales over time, averaged across all states, regions, and sectors. The data show pronounced seasonal patterns with consistent annual cycles and relatively stable long-run average levels.

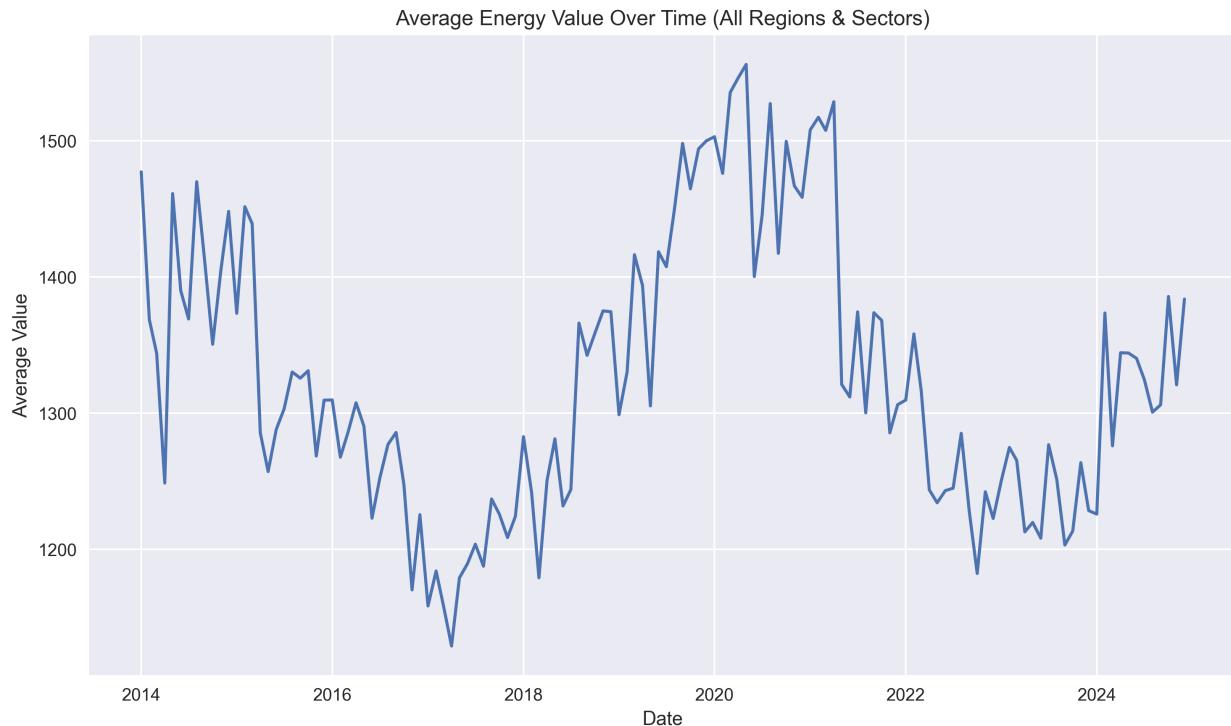


Figure 1: Average retail electricity sales over time (all regions and sectors).

3.3 Distribution by Region and Sector

Figure 2 summarizes the distribution of sales by region and sector. The South region shows the highest consumption levels, while the Transportation sector exhibits minimal electricity usage relative to other sectors.

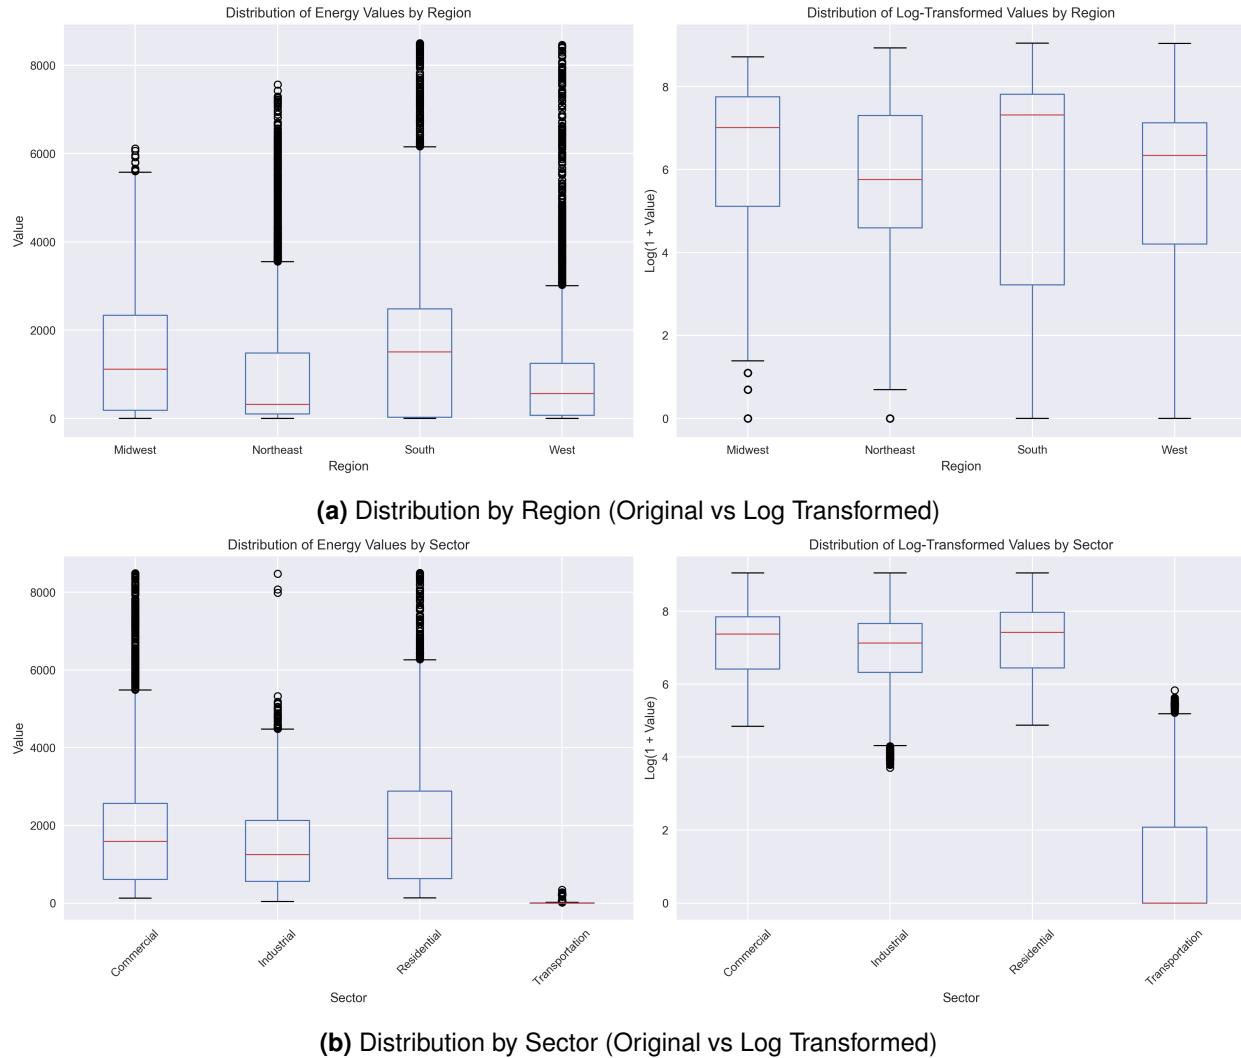


Figure 2: Boxplots of retail electricity sales distributions.

3.4 Normality Assessment

Figure 3 compares the Quantile-Quantile (QQ) plots of the original and log-transformed data. The transformation significantly improves the normality of the residuals, particularly by compressing the extreme right tail of the distribution.

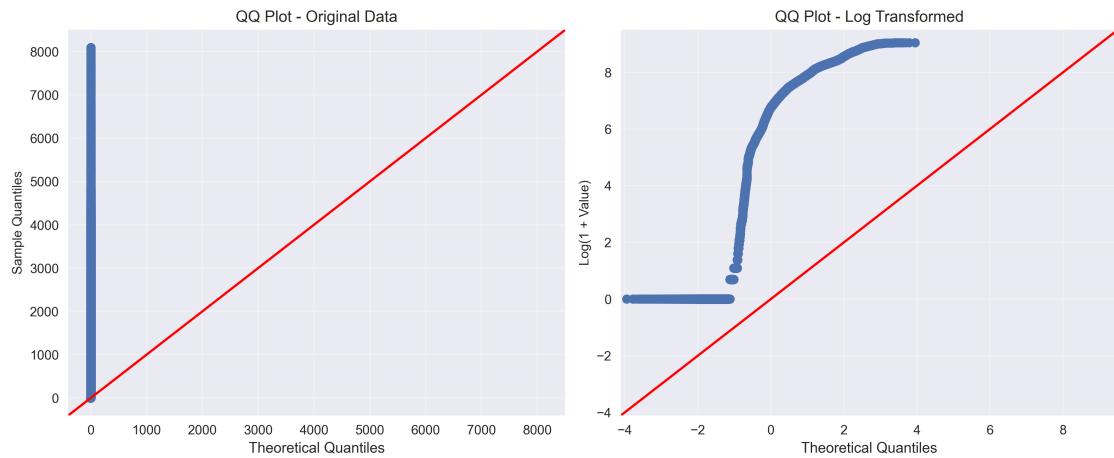


Figure 3: QQ plots showing improvement in normality after log transformation.

3.5 Regional Time Series

Figures 4 through 6 break down time series trends by geography. The South (Figure 4) exhibits the highest aggregate demand. The seasonal breakdowns for the Midwest and Northeast (Figures 5 and 6) illustrate distinct winter and summer peaks driven by heating and cooling loads.

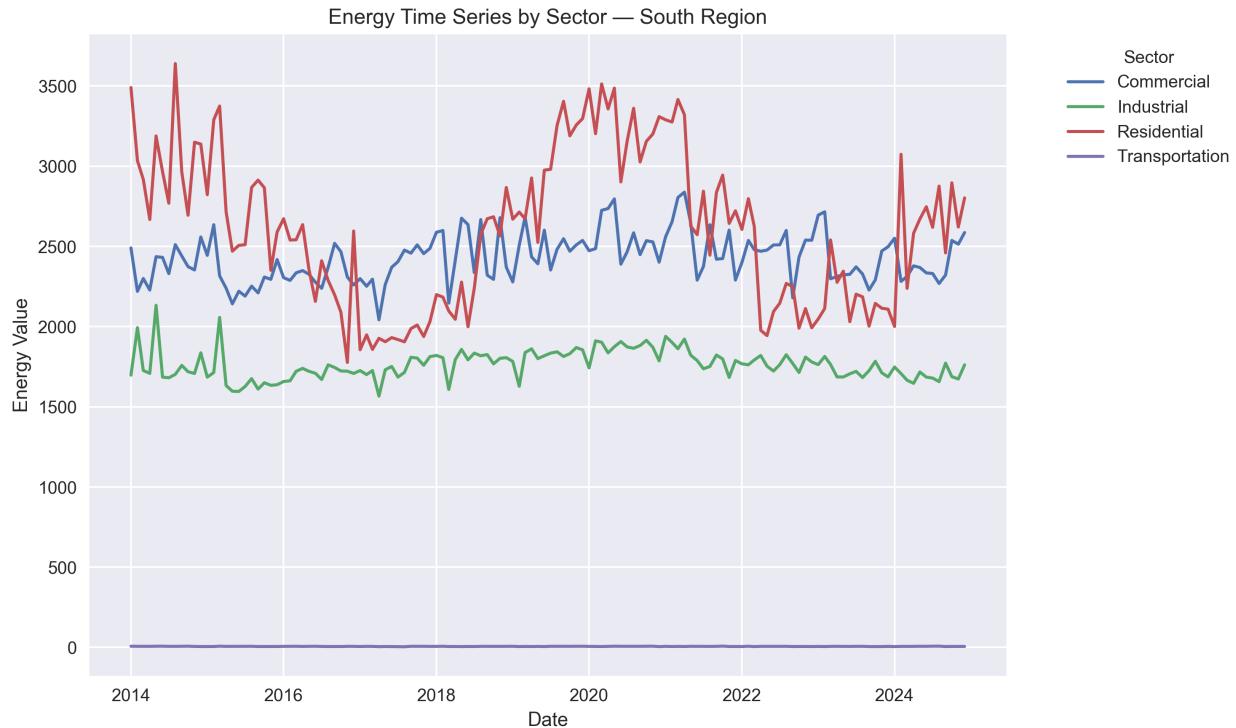


Figure 4: South Region: Time series by sector showing dominance of residential demand.



Figure 5: Midwest Region: Seasonal decomposition of electricity sales.

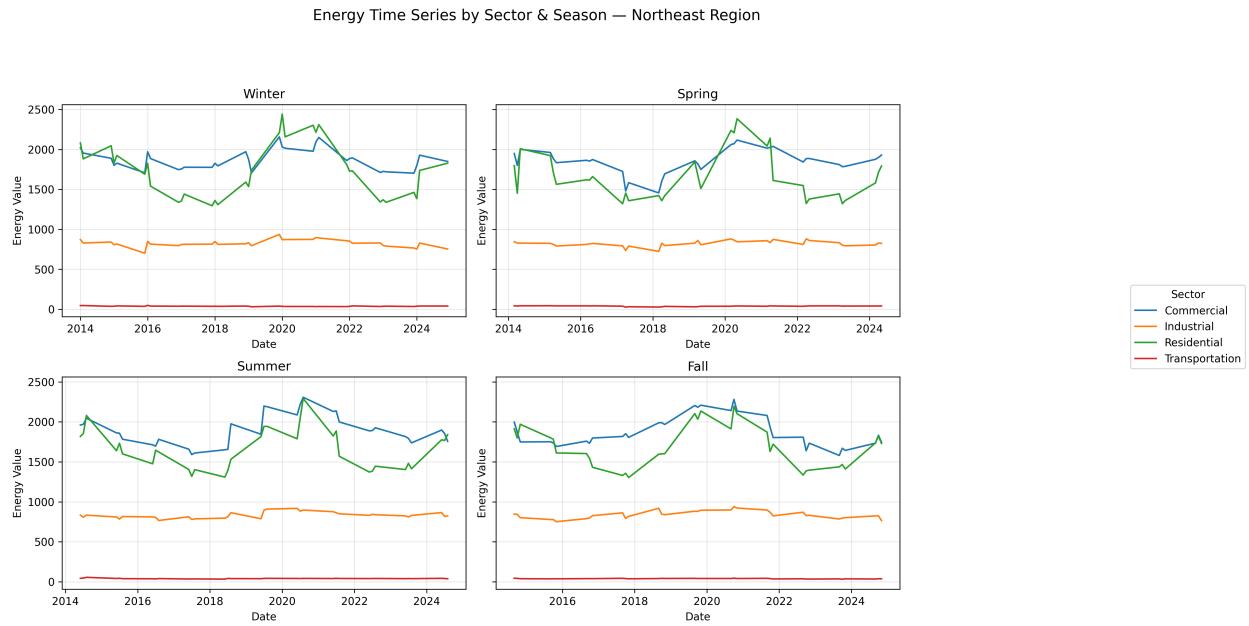


Figure 6: Northeast Region: Seasonal decomposition of electricity sales.

3.6 Sector-Specific Normality

Figure 7 displays the distributional characteristics of each sector. The Transportation sector shows the most deviation from normality due to the high frequency of near-zero values.

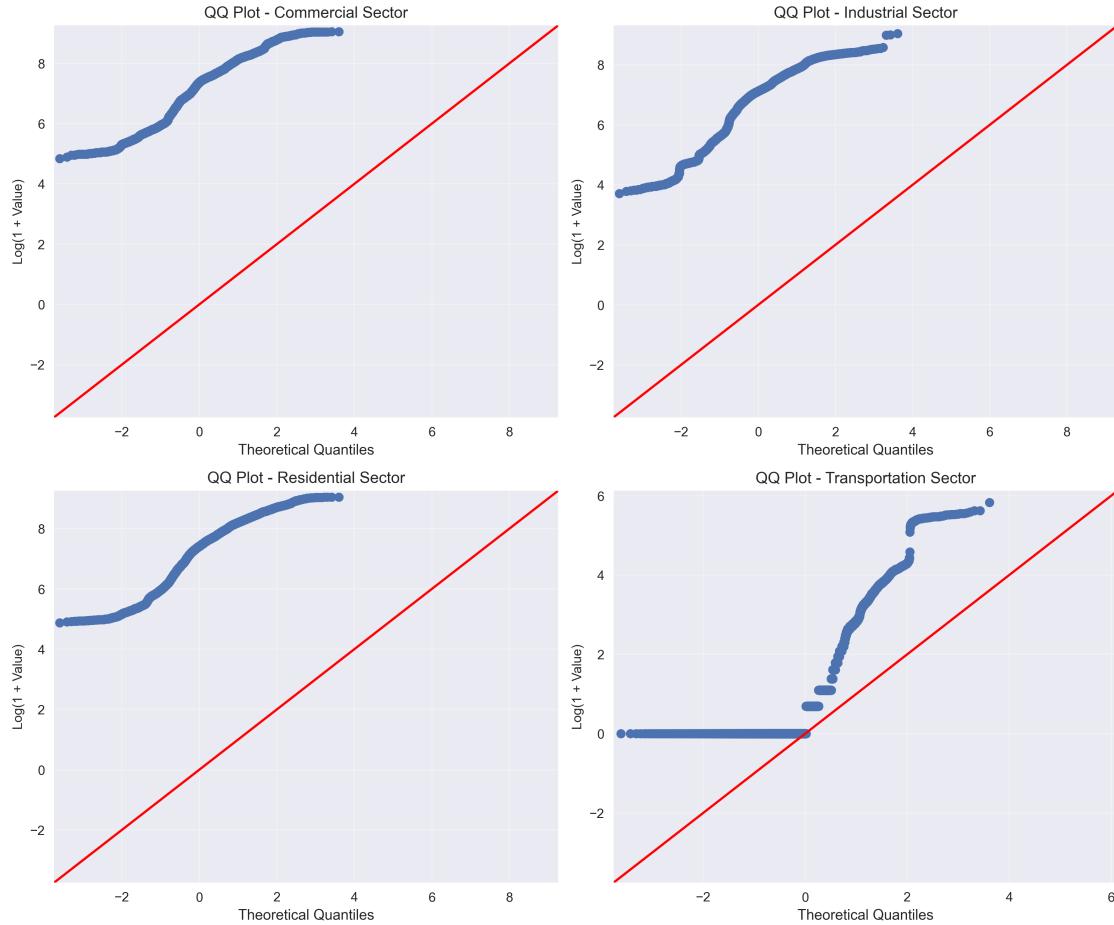


Figure 7: QQ plots by economic sector (Log-Transformed).

4 Grouped Summary Statistics

Table 2 details electricity consumption broken down by region and sector. The **South** region consistently shows the highest mean consumption across Residential and Commercial sectors.

Table 2: Retail electricity sales (MWh) by region and sector.

Region	Sector	Count	Mean	Std. Dev.	Min	Max
Midwest	Commercial	1,584	1,960.33	1,190.33	332	4,968
	Industrial	1,584	1,982.94	1,226.24	211	4,608
	Residential	1,584	2,037.79	1,321.82	287	6,110
	Transportation	1,584	4.18	11.56	0	79
Northeast	Commercial	1,187	1,872.30	1,925.55	126	7,559
	Industrial	1,188	829.36	1,220.17	40	4,693
	Residential	1,188	1,684.84	1,666.40	141	6,318
	Transportation	1,188	38.71	69.47	0	341
South	Commercial	1,941	2,427.71	1,731.01	255	8,489
	Industrial	1,983	1,762.83	868.89	92	8,481
	Residential	1,863	2,607.98	1,506.32	271	8,499
	Transportation	2,112	5.79	10.93	0	68
West	Commercial	1,603	1,133.30	1,086.34	180	8,419
	Industrial	1,716	1,128.44	979.19	94	5,327
	Residential	1,674	1,488.26	1,694.28	130	8,463
	Transportation	1,716	6.27	16.38	0	83

5 Statistical Analysis

5.1 Regression Model Results

A multiple linear regression model was fit to the log-transformed data ($\ln(\text{Value})$). The model demonstrated a strong fit ($R^2 = 0.852$), suggesting that region and sector are highly effective predictors of electricity consumption.

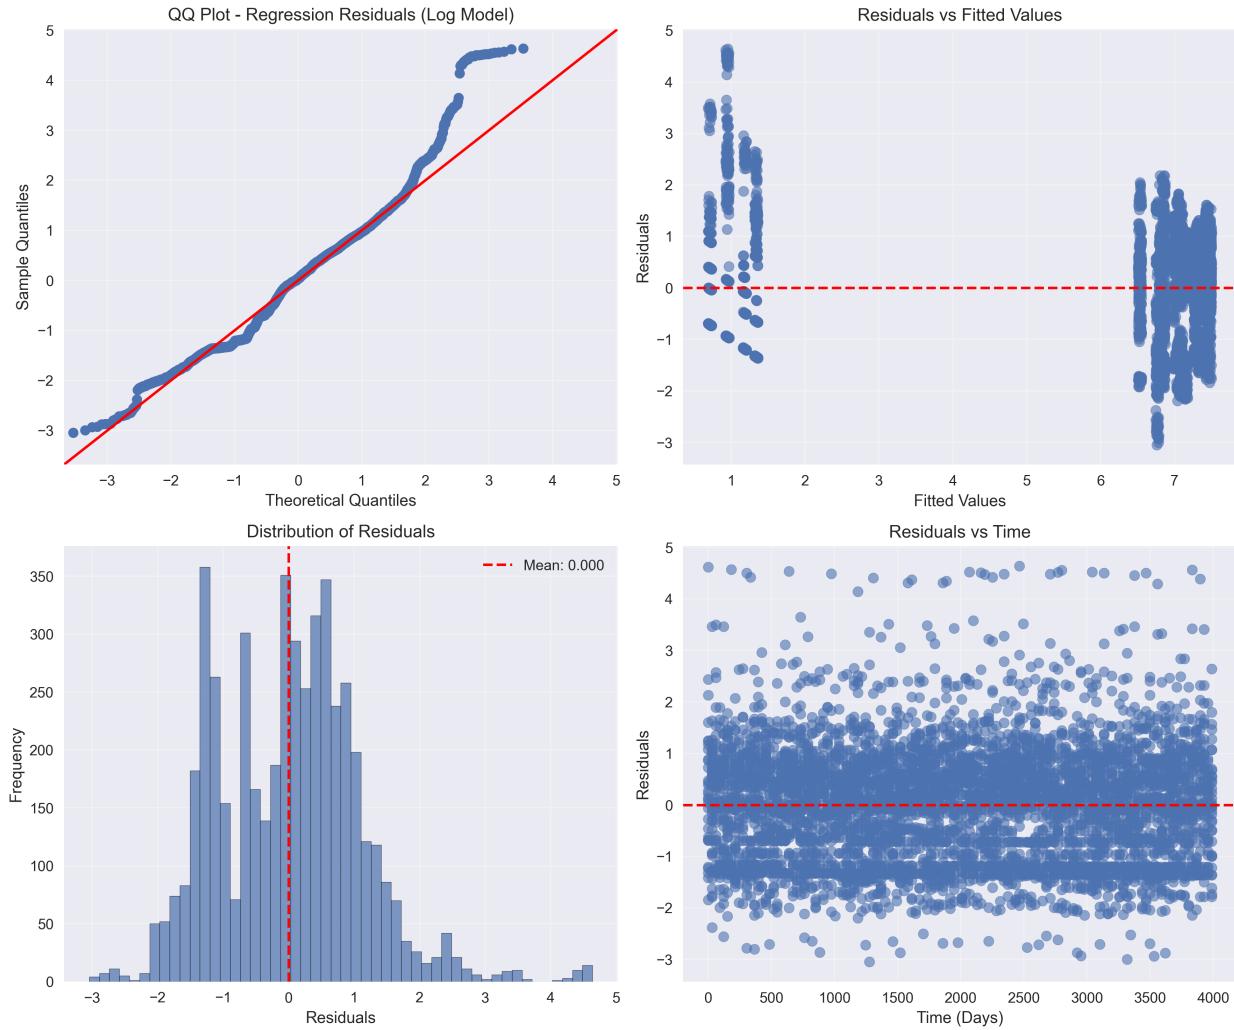


Figure 8: Regression diagnostics: The Residuals vs Fitted plot (top right) and Residuals vs Time (bottom right) show homoscedasticity, validating the log-linear specification.

5.2 Interpretation

- **Regional Effects:** The **South** region exhibits significantly higher consumption relative to the Midwest baseline ($\sim +58\%$), attributable to climate-driven cooling demand.
- **Sectoral Effects:** The **Transportation** sector shows a massive negative coefficient relative to Commercial ($\sim 99.8\%$ lower), confirming that electric transportation grid load was negligible during the 2014-2024 period compared to building stock.
- **Time Trend:** No statistically significant linear time trend was found after controlling for structural factors.

6 Discussion

This analysis provides a comprehensive statistical evaluation of U.S. retail electricity sales from 2014 to 2024. The defining characteristic of the dataset is the dominance of structural factors over temporal ones. With a model R^2 of approximately 0.852, the combination of geographic region and economic sector explains the vast majority of variance in electricity consumption. This high degree of explanatory power suggests that U.S. electricity demand is relatively deterministic based on location and activity type, rather than being driven by a uniform national growth trend.

6.1 The “South Effect” and Regional Heterogeneity

The most pronounced finding in the spatial analysis is the consumption premium associated with the **South** census region. Controlling for sector and time, the South exhibits significantly higher electricity usage than the Midwest, Northeast, or West. This disparity is attributable to a convergence of climatological and economic drivers:

- **Cooling Degree Days (CDD):** The South experiences a higher frequency of extreme heat events compared to other regions. Unlike heating, which can be fueled by natural gas or oil, cooling is almost exclusively electric. This creates a high “cooling penalty” that raises baseload demand.
- **Electrification of Heating:** In many parts of the Southeast, electric resistance heating and heat pumps are more common than in the Northeast (where fuel oil and gas dominate), leading to a dual-peak load profile (high summer and moderate winter demand).
- **Economic Geography:** The region has seen faster population growth and industrial expansion (e.g., petrochemicals, auto manufacturing) than the Rust Belt, increasing aggregate load in absolute terms.

In contrast, the **West** displays lower consumption relative to its economic output. This phenomenon validates the long-term impact of aggressive energy efficiency policies (such as California’s Title 24) and a milder coastal climate that suppresses total HVAC demand.

6.2 The Transportation Gap: A Looming Grid Shock

The statistical results for the **Transportation** sector are stark. With a regression coefficient corresponding to $\sim 99.8\%$ lower consumption than the Commercial sector, transportation electricity sales are effectively negligible in the context of the total grid load during the 2014–2024 period.

Implication: This near-zero baseline serves as a critical warning for infrastructure planning. Despite the cultural visibility of electric vehicles (EVs), their aggregate demand has not yet structurally altered the load profile of the U.S. grid. However, as the vehicle fleet transitions from liquid fuels to electricity, the grid faces a massive latent demand shock. If the Transportation sector were to grow to even 20% of the Residential sector's consumption volume, it would require a historic expansion of generation and transmission capacity. The current stability of the grid is partly due to this transition being in its infancy.

6.3 Economic Decoupling: The Absence of a Time Trend

Perhaps the most counter-intuitive finding is the lack of a statistically significant linear time trend ($p > 0.05$) in the regression model. Given that the U.S. population and GDP grew between 2014 and 2024, the expectation might be a corresponding linear increase in electricity sales.

The flatness of the time dimension suggests a phenomenon known as **Energy-GDP Decoupling**:

1. **Efficiency Gains:** Improvements in end-use efficiency (LED lighting, higher SEER ratings for AC units, industrial process optimization) have likely negated the load growth from new housing and devices.
2. **Structural Shift:** The U.S. economy continues to pivot toward service-oriented sectors (Commercial) and away from heavy manufacturing (Industrial), reducing the energy intensity per dollar of GDP.
3. **Behind-the-Meter Generation:** The growth of residential solar reduces *net* retail sales (the variable measured here) without necessarily reducing actual consumption, masking the true demand curve.

6.4 Methodological Insights: The Power of Log-Normality

The dramatic improvement in model fit when moving from raw values ($R^2 \approx 0.24$) to log-transformed values ($R^2 \approx 0.85$) is not merely a statistical technicality—it reveals the physics of energy usage.

Electricity consumption data is **heteroscedastic** and **heavy-tailed**. Large consumers (industrial plants, large commercial buildings) vary in their usage by orders of magnitude more than small consumers (households). Furthermore, effects are multiplicative: an extreme heatwave does not add a fixed number of kilowatt-hours to every building; rather, it multiplies the cooling load by a factor relative to the building's size. The log-linear model correctly captures this multiplicative dynamic, providing a robust framework for

future forecasting.

6.5 Strategic Recommendations

Based on these findings, energy planners should prioritize:

- **Regional Capacity Markets:** The South's distinct load profile requires specific capacity market mechanisms to manage high cooling peaks, distinct from the heating-focused needs of the Northeast.
- **EV Infrastructure Readiness:** Utilities must prepare for the Transportation sector to move from a statistical error to a dominant load class. This requires upgrading distribution transformers and incentivizing off-peak charging before the load fully materializes.
- **Efficiency as a Resource:** The stable time trend proves that efficiency programs are effective at offsetting growth. Continued investment in efficiency is the most cost-effective method to maintain grid reliability in the face of the coming electrification wave.

7 Conclusion

This analysis demonstrates that understanding U.S. electricity consumption requires recognizing its fundamentally structural nature. The strong regional and sectoral patterns identified here suggest that effective energy policy must be geographically differentiated and sectorally targeted. The success of the log transformation in capturing the underlying data structure provides both a methodological lesson for energy analysts and substantive insight into the multiplicative nature of consumption drivers.

As the energy transition accelerates with transportation electrification, renewable energy integration, and climate-driven demand changes, the structural patterns identified in this analysis will provide a crucial baseline for understanding how these transformations reshape America's electricity landscape. The regional and sectoral differences documented here represent not just statistical patterns but fundamental features of the American economy, climate, and society that will continue to shape energy outcomes for decades to come.

References

- [1] Hunter, J. D. (2007). Matplotlib: A 2D graphics environment. *Computing in Science & Engineering*, 9(3), 90–95. <https://doi.org/10.1109/MCSE.2007.55>
- [2] McKinney, W. (2010). Data structures for statistical computing in Python. In S. van der Walt & J. Millman (Eds.), *Proceedings of the 9th Python in Science Conference* (pp. 51–56).
- [3] Seabold, S., & Perktold, J. (2010). Statsmodels: Econometric and statistical modeling with Python. In *Proceedings of the 9th Python in Science Conference* (pp. 92–96).
- [4] U.S. Energy Information Administration. (2024). *Electric power monthly: Retail sales of electricity to ultimate customers* [Data set]. U.S. Department of Energy. <https://www.eia.gov/electricity/monthly/>
- [5] Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D., ... & van der Walt, S. J. (2020). SciPy 1.0: Fundamental algorithms for scientific computing in Python. *Nature Methods*, 17, 261–272. <https://doi.org/10.1038/s41592-019-0686-2>

A Python Source Code

The following Python script was used to perform data cleaning, exploratory data analysis, and statistical modeling as described in Section 2.

```

1 #!/usr/bin/env python
2 # coding: utf-8
3
4 # =====
5 # Time Series Energy Data: Analysis Notebook
6 # =====
7
8 import pandas as pd
9 import numpy as np
10 import matplotlib.pyplot as plt
11 from pathlib import Path
12 from IPython.display import display
13
14 import statsmodels.api as sm
15 import statsmodels.formula.api as smf
16 from statsmodels.stats.anova import anova_lm
17 from scipy import stats
18
19 # Make plots a bit nicer
20 plt.rcParams["figure.figsize"] = (10, 6)
21 plt.rcParams["figure.dpi"] = 110
22
23 # 1. ----- CONFIG -----
24
25 # Path to your CSV file
26 DATA_PATH = Path(
27     r"C:\Users\mokol\Desktop\Past FIU Classes\CAP5768 Intro to Data
28     Science\Retail Sales of Electricity by State - CLEANED.csv"
29 )
30
31 # Output directory for figures
32 FIGURES_PATH = Path(r"C:\Users\mokol\Desktop\Current FIU Classes\STA6244
33     Data Analysis I\Final Project\figures")
34 FIGURES_PATH.mkdir(parents=True, exist_ok=True) # Create directory if it
35     doesn't exist
36
37 # Name of the column that contains state-sector labels, like "AL-
38     Residential"
39 STATE_SECTOR_COL = "State_Sector"
```

```

37 # 2. ----- STATE REGION MAPPING -----
38
39 state_to_region = {
40     # Northeast
41     "ME": "Northeast", "NH": "Northeast", "VT": "Northeast", "MA": "Northeast",
42     "RI": "Northeast", "CT": "Northeast", "NY": "Northeast", "NJ": "Northeast",
43     "PA": "Northeast",
44
45     # Midwest
46     "OH": "Midwest", "IN": "Midwest", "IL": "Midwest", "MI": "Midwest",
47     "WI": "Midwest", "MN": "Midwest", "IA": "Midwest", "MO": "Midwest",
48     "ND": "Midwest", "SD": "Midwest", "NE": "Midwest", "KS": "Midwest",
49
50     # South
51     "DE": "South", "MD": "South", "DC": "South", "VA": "South",
52     "WV": "South", "NC": "South", "SC": "South", "GA": "South",
53     "FL": "South", "KY": "South", "TN": "South", "AL": "South",
54     "MS": "South", "AR": "South", "LA": "South", "OK": "South",
55     "TX": "South",
56
57     # West
58     "MT": "West", "ID": "West", "WY": "West", "CO": "West",
59     "NM": "West", "AZ": "West", "UT": "West", "NV": "West",
60     "WA": "West", "OR": "West", "CA": "West", "AK": "West",
61     "HI": "West",
62 }
63
64 # 3. ----- LOAD THE DATA -----
65
66 df = pd.read_csv(DATA_PATH)
67
68 # Ensure the state-sector column exists
69 if STATE_SECTOR_COL not in df.columns:
70     raise ValueError(
71         f"Column '{STATE_SECTOR_COL}' not found in CSV columns: {df.
72         columns.tolist()}")
73
74 print("Raw wide data (first 5 rows):")
75 display(df.head())
76
77 # 4. ----- SPLIT STATE-SECTOR COLUMN -----
78

```

```

79 # We assume the format is "<STATE>-<SECTOR>", e.g. "AL-Residential"
80 state_sector_split = df[STATE_SECTOR_COL].astype(str).str.split("-", n=1,
81     expand=True)
82 df["State"] = state_sector_split[0].str.strip()
83 df["Sector"] = state_sector_split[1].str.strip()
84
85 # Remove all region-level totals (Sector == "Total")
86 df = df[df["Sector"].str.lower() != "total"].copy()
87
88 print("\nAfter dropping Sector == 'Total':")
89 display(df.head())
90
91
92 # 5. ----- IDENTIFY DATE COLUMNS -----
93 id_cols = [STATE_SECTOR_COL, "State", "Sector"]
94 date_cols = [c for c in df.columns if c not in id_cols]
95
96 print("\nDetected date columns (truncated):", date_cols[:10], "...")
97 print("Number of date columns:", len(date_cols))
98
99 # Parse date headers like "2014 January" using explicit format
100 parsed_dates = pd.to_datetime(date_cols, format="%Y %B", errors="coerce")
101
102 if parsed_dates.isnull().all():
103     print("\nWarning: Could not parse date column names as datetimes.
104           Using them as strings.")
105 else:
106     new_date_cols = {old: new for old, new in zip(date_cols, parsed_dates)}
107     df = df.rename(columns=new_date_cols)
108     date_cols = list(new_date_cols.values())
109     print("\nSuccessfully parsed date column names as datetimes.")
110
111 # 6. ----- WIDE      LONG (TIDY) FORMAT -----
112
113 df_long = df.melt(
114     id_vars=["State", "Sector"],
115     value_vars=date_cols,
116     var_name="Date",
117     value_name="Value"
118 )
119
120 # Convert Date to datetime if possible
121 df_long["Date"] = pd.to_datetime(df_long["Date"], errors="coerce")

```

```
121 # --- ADD SEASON FEATURE (WINTER / SPRING / SUMMER / FALL) ---
122
123 # Drop rows with invalid dates first (if any)
124 df_long = df_long.dropna(subset=["Date"]).copy()
125
126 # Extract month
127 df_long["Month"] = df_long["Date"].dt.month
128
129 def month_to_season(m):
130     if m in [12, 1, 2]:
131         return "Winter"
132     elif m in [3, 4, 5]:
133         return "Spring"
134     elif m in [6, 7, 8]:
135         return "Summer"
136     else: # 9, 10, 11
137         return "Fall"
138
139 df_long["Season"] = df_long["Month"].apply(month_to_season)
140
141
142 print("\nLong/tidy data before cleaning Value (first 5 rows):")
143 display(df_long.head())
144
145 # Ensure Value is numeric
146 df_long["Value"] = (
147     df_long["Value"]
148     .astype(str)
149     .str.replace(", ", "", regex=False)
150     .str.strip()
151 )
152 df_long["Value"] = pd.to_numeric(df_long["Value"], errors="coerce")
153
154 # Drop rows with missing Value
155 df_long = df_long.dropna(subset=["Value"])
156
157 print("\nLong/tidy data after cleaning Value (first 5 rows):")
158 display(df_long.head())
159 print("Value dtype:", df_long["Value"].dtype)
160
161 # 7. ----- MAP STATES TO REGIONS -----
162
163 df_long["Region"] = df_long["State"].map(state_to_region)
164
165 missing_region_mask = df_long["Region"].isna()
```

```

166 if missing_region_mask.any():
167     missing_states = df_long.loc[missing_region_mask, "State"].unique()
168     print("\nWarning: These states had no region mapping and will be
169     dropped:", missing_states)
170     df_long = df_long.loc[~missing_region_mask].copy()
171
172 print("\nTidy data after region mapping (first 5 rows):")
173 display(df_long.head())
174
175 print("\nNumber of observations:", len(df_long))
176 print("Columns:", df_long.columns.tolist())
177
178 # 8. ---- DATA FILTERING AND CLEANING ----
179
180 print("\n===== DATA FILTERING AND CLEANING =====")
181
182 # Remove extreme outliers using IQR method
183 Q1 = df_long["Value"].quantile(0.25)
184 Q3 = df_long["Value"].quantile(0.75)
185 IQR = Q3 - Q1
186 lower_bound = Q1 - 3 * IQR # Using 3*IQR for less aggressive filtering
187 upper_bound = Q3 + 3 * IQR
188
189 original_count = len(df_long)
190 df_long = df_long[(df_long["Value"] >= lower_bound) & (df_long["Value"] <=
191     upper_bound)].copy()
192 filtered_count = len(df_long)
193
194 print(f"Removed {original_count - filtered_count} extreme outliers ({(
195     original_count - filtered_count)/original_count}*100:.1f}% of data)")
196 print(f"Final dataset size: {filtered_count} observations")
197
198 # Create log-transformed variable for analysis
199 df_long["Log_Value"] = np.log1p(df_long["Value"])
200
201 # 9. ---- AGGREGATE BY REGION + SECTOR + DATE --
202
203 region_sector_ts = (
204     df_long
205     .groupby(["Region", "Sector", "Date", "Season"], as_index=False)[
206         "Value"]
207     .mean() # change to .sum() if you prefer total usage per region
208 )
209
210 print("\nAggregated time series (Region, Sector, Date):")

```

```
207 display(region_sector_ts.head())
208
209 # 10. ----- PLOT PER REGION (BY SEASON) -----
210 # Each figure = one region
211 # Inside each figure: 4 subplots (Winter, Spring, Summer, Fall)
212 # Each line in a subplot = one sector
213
214 regions = sorted(region_sector_ts["Region"].unique())
215 season_order = ["Winter", "Spring", "Summer", "Fall"]
216
217 for region in regions:
218     region_df = region_sector_ts[region_sector_ts["Region"] == region].copy()
219
220     fig, axes = plt.subplots(2, 2, figsize=(14, 8), sharey=True)
221     axes = axes.flatten()
222
223     for ax, season in zip(axes, season_order):
224         season_df = region_df[region_df["Season"] == season].copy()
225         if season_df.empty:
226             ax.set_title(f"{season} (no data)")
227             ax.axis("off")
228             continue
229
230         pivot = season_df.pivot(
231             index="Date",
232             columns="Sector",
233             values="Value"
234         ).sort_index()
235
236         for sector in pivot.columns:
237             ax.plot(pivot.index, pivot[sector], label=sector)
238
239         ax.set_title(season)
240         ax.set_xlabel("Date")
241         ax.set_ylabel("Energy Value")
242         ax.grid(True, alpha=0.3)
243
244     # Put one legend for the whole figure on the right
245     handles, labels = axes[0].get_legend_handles_labels()
246     fig.legend(handles, labels, title="Sector",
247                bbox_to_anchor=(1.05, 0.5), loc="center left")
248
249 fig.suptitle(f"Energy Time Series by Sector & Season      {region}
Region",
```

```
250         fontsize=14, y=1.02)
251     plt.tight_layout(rect=[0, 0, 0.85, 0.95])
252
253     # Save the figure
254     filename = FIGURES_PATH / f"timeseries_{region.lower()}_by_season.png"
255     plt.savefig(filename, dpi=300, bbox_inches='tight')
256     print(f"Saved: {filename}")
257
258     plt.show()
259
260
261
262 # =====
263 # PART 2: DESCRIBE THE DATASET (SUMMARY & VISUALIZATIONS)
264 # =====
265
266 print("\n===== BASIC SUMMARY STATISTICS (NUMERIC) =====")
267 print("Original Values:")
268 display(df_long[["Value"]].describe())
269 print("\nLog-Transformed Values:")
270 display(df_long[["Log_Value"]].describe())
271
272 print("\n===== CATEGORY COUNTS =====")
273 print("\nSector counts:")
274 display(df_long[["Sector"]].value_counts())
275 print("\nRegion counts:")
276 display(df_long[["Region"]].value_counts())
277
278 # Grouped summary by Sector and Region
279 group_summary = (
280     df_long
281     .groupby(["Region", "Sector"])["Value"]
282     .agg(["count", "mean", "std", "min", "max"])
283     .reset_index()
284     .sort_values(["Region", "Sector"])
285 )
286
287 print("\n===== GROUPED SUMMARY BY REGION & SECTOR =====")
288 display(group_summary)
289
290 # ----- VISUALIZATIONS -----
291
292 # 1) Distribution of Values by Sector (boxplot)
293 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))
```

```
295 df_long.boxplot(column="Value", by="Sector", rot=45, ax=ax1)
296 ax1.set_title("Distribution of Energy Values by Sector")
297 ax1.set_xlabel("Sector")
298 ax1.set_ylabel("Value")
299
300 df_long.boxplot(column="Log_Value", by="Sector", rot=45, ax=ax2)
301 ax2.set_title("Distribution of Log-Transformed Values by Sector")
302 ax2.set_xlabel("Sector")
303 ax2.set_ylabel("Log(1 + Value)")
304
305 plt.suptitle("")
306 plt.tight_layout()
307
308 # Save the figure
309 filename = FIGURES_PATH / "boxplot_sector_comparison.png"
310 plt.savefig(filename, dpi=300, bbox_inches='tight')
311 print(f"Saved: {filename}")
312
313 plt.show()
314
315 # 2) Distribution of Values by Region (boxplot)
316 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))
317
318 df_long.boxplot(column="Value", by="Region", rot=0, ax=ax1)
319 ax1.set_title("Distribution of Energy Values by Region")
320 ax1.set_xlabel("Region")
321 ax1.set_ylabel("Value")
322
323 df_long.boxplot(column="Log_Value", by="Region", rot=0, ax=ax2)
324 ax2.set_title("Distribution of Log-Transformed Values by Region")
325 ax2.set_xlabel("Region")
326 ax2.set_ylabel("Log(1 + Value)")
327
328 plt.suptitle("")
329 plt.tight_layout()
330
331 # Save the figure
332 filename = FIGURES_PATH / "boxplot_region_comparison.png"
333 plt.savefig(filename, dpi=300, bbox_inches='tight')
334 print(f"Saved: {filename}")
335
336 plt.show()
337
338 # 3) Average Value over Time (overall)
339 avg_over_time = (
```

```
340 df_long
341 .groupby("Date", as_index=True) ["Value"]
342 .mean()
343 .sort_index()
344 )
345
346 plt.figure()
347 plt.plot(avg_over_time.index, avg_over_time.values)
348 plt.title("Average Energy Value Over Time (All Regions & Sectors)")
349 plt.xlabel("Date")
350 plt.ylabel("Average Value")
351 plt.grid(True)
352 plt.tight_layout()
353
354 # Save the figure
355 filename = FIGURES_PATH / "timeseries_overall_average.png"
356 plt.savefig(filename, dpi=300, bbox_inches='tight')
357 print(f"Saved: {filename}")
358
359 plt.show()
360
361
362 # =====
363 # PART 3: QQ PLOTS FOR NORMALITY CHECK
364 # =====
365
366 print("\n===== QQ PLOTS FOR NORMALITY ANALYSIS =====")
367
368 def create_qqplot(data, title, ax, transform=False):
369     """Create QQ plot with proper formatting"""
370     if transform:
371         plot_data = np.log1p(data)
372         ylabel = "Log(1 + Value)"
373     else:
374         plot_data = data
375         ylabel = "Sample Quantiles"
376
377     # Remove extreme outliers for better visualization
378     Q1 = np.percentile(plot_data, 25)
379     Q3 = np.percentile(plot_data, 75)
380     IQR = Q3 - Q1
381     lower_bound = Q1 - 3 * IQR
382     upper_bound = Q3 + 3 * IQR
383
384     filtered_data = plot_data[(plot_data >= lower_bound) & (plot_data <=
```

```
    upper_bound)]  
385  
386     sm.qqplot(filtered_data, line='45', ax=ax)  
387     ax.set_title(f"{title}", fontsize=12)  
388     ax.set_ylabel(ylabel, fontsize=10)  
389     ax.grid(True, alpha=0.3)  
390  
391     return ax  
392  
393 # 1. QQ Plots by Region  
394 print("\n--- QQ Plots by Region ---")  
395 regions = sorted(df_long["Region"].unique())  
396 fig, axes = plt.subplots(2, 2, figsize=(12, 10))  
397 axes = axes.flatten()  
398  
399 for i, region in enumerate(regions):  
400     if i < len(axes):  
401         subset = df_long[df_long["Region"] == region][["Value"]]  
402         create_qqplot(subset, f"QQ Plot - {region} Region", axes[i],  
403                      transform=True)  
404  
405 plt.tight_layout()  
406  
407 # Save the figure  
408 filename = FIGURES_PATH / "qqplot_by_region.png"  
409 plt.savefig(filename, dpi=300, bbox_inches='tight')  
410 print(f"Saved: {filename}")  
411 plt.show()  
412  
413 # 2. QQ Plots by Sector  
414 print("\n--- QQ Plots by Sector ---")  
415 sectors = sorted(df_long["Sector"].unique())  
416 fig, axes = plt.subplots(2, 2, figsize=(12, 10))  
417 axes = axes.flatten()  
418  
419 for i, sector in enumerate(sectors):  
420     if i < len(axes):  
421         subset = df_long[df_long["Sector"] == sector][["Value"]]  
422         create_qqplot(subset, f"QQ Plot - {sector} Sector", axes[i],  
423                      transform=True)  
424  
425 plt.tight_layout()  
426 # Save the figure
```

```
427 filename = FIGURES_PATH / "qqplot_by_sector.png"
428 plt.savefig(filename, dpi=300, bbox_inches='tight')
429 print(f"Saved: {filename}")
430
431 plt.show()
432
433 # 3. Overall QQ Plots
434 print("\n--- Overall QQ Plots ---")
435 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))
436
437 create_qqplot(df_long["Value"], "QQ Plot - Original Data", ax1, transform=False)
438 create_qqplot(df_long["Value"], "QQ Plot - Log Transformed", ax2, transform=True)
439
440 plt.tight_layout()
441
442 # Save the figure
443 filename = FIGURES_PATH / "qqplot_overall_comparison.png"
444 plt.savefig(filename, dpi=300, bbox_inches='tight')
445 print(f"Saved: {filename}")
446
447 plt.show()
448
449
450 # =====
451 # PART 4: STATISTICAL ANALYSIS (REGRESSION & ANOVA)
452 # =====
453
454 # Create a numeric time variable (days since earliest date)
455 if df_long["Date"].notna().any():
456     min_date = df_long["Date"].min()
457     df_long["TimeNumeric"] = (df_long["Date"] - min_date).dt.days
458 else:
459     print("\nDate could not be parsed; using simple index as time.")
460     df_long = df_long.sort_values(["State", "Sector"])
461     df_long["TimeNumeric"] = df_long.groupby(["State", "Sector"]).cumcount()
462
463 print("\nTimeNumeric summary:")
464 display(df_long["TimeNumeric"].describe())
465
466 # Sample for regression (using filtered data)
467 sample_size = min(5000, len(df_long))
468 df_sample = df_long.sample(n=sample_size, random_state=42).copy()
```

```
469
470 print(f"\nUsing {len(df_sample)} rows for regression sample.")
471
472 # Multiple Linear Regression with LOG TRANSFORMED data
473 print("\n===== MULTIPLE LINEAR REGRESSION (Log-Transformed) =====")
474 formula_log = "Log_Value ~ TimeNumeric + C(Region) + C(Sector)"
475 model_log = smf.ols(formula=formula_log, data=df_sample).fit()
476
477 print(model_log.summary())
478
479 # ANOVA to test Region & Sector effects
480 anova_results_log = anova_lm(model_log, typ=2)
481 print("\n===== ANOVA TABLE (Type II) - Log Transformed =====")
482 print(anova_results_log)
483
484 # Also run on original data for comparison
485 print("\n===== MULTIPLE LINEAR REGRESSION (Original Data) =====")
486 formula_orig = "Value ~ TimeNumeric + C(Region) + C(Sector)"
487 model_orig = smf.ols(formula=formula_orig, data=df_sample).fit()
488
489 print(model_orig.summary())
490
491 # =====
492 # PART 5: REGRESSION DIAGNOSTICS
493 # =====
494
495 print("\n===== REGRESSION DIAGNOSTICS =====")
496
497 # Residuals analysis for log-transformed model
498 fig, axes = plt.subplots(2, 2, figsize=(12, 10))
499
500 # 1. QQ Plot of residuals
501 sm.qqplot(model_log.resid, line='45', ax=axes[0,0])
502 axes[0,0].set_title('QQ Plot - Regression Residuals (Log Model)')
503 axes[0,0].grid(True, alpha=0.3)
504
505 # 2. Residuals vs Fitted
506 axes[0,1].scatter(model_log.fittedvalues, model_log.resid, alpha=0.6)
507 axes[0,1].axhline(y=0, color='red', linestyle='--')
508 axes[0,1].set_xlabel('Fitted Values')
509 axes[0,1].set_ylabel('Residuals')
510 axes[0,1].set_title('Residuals vs Fitted Values')
511 axes[0,1].grid(True, alpha=0.3)
512
513 # 3. Distribution of residuals
```

```
514 axes[1,0].hist(model_log.resid, bins=50, alpha=0.7, edgecolor='black')
515 axes[1,0].axvline(model_log.resid.mean(), color='red', linestyle='--',
516                     label=f'Mean: {model_log.resid.mean():.3f}')
517 axes[1,0].set_xlabel('Residuals')
518 axes[1,0].set_ylabel('Frequency')
519 axes[1,0].set_title('Distribution of Residuals')
520 axes[1,0].legend()
521 axes[1,0].grid(True, alpha=0.3)
522
523 # 4. Residuals over time
524 axes[1,1].scatter(df_sample["TimeNumeric"], model_log.resid, alpha=0.6)
525 axes[1,1].axhline(y=0, color='red', linestyle='--')
526 axes[1,1].set_xlabel('Time (Days)')
527 axes[1,1].set_ylabel('Residuals')
528 axes[1,1].set_title('Residuals vs Time')
529 axes[1,1].grid(True, alpha=0.3)
530
531 plt.tight_layout()
532
533 # Save the figure
534 filename = FIGURES_PATH / "regression_diagnostics.png"
535 plt.savefig(filename, dpi=300, bbox_inches='tight')
536 print(f"Saved: {filename}")
537
538 plt.show()
539
540 # =====
541 # PART 6: BASIC INTERPRETATION HELPERS
542 # =====
543
544 r_squared_log = model_log.rsquared
545 adj_r_squared_log = model_log.rsquared_adj
546 coeffs_log = model_log.params
547
548 print("\n===== QUICK MODEL INTERPRETATION (Log Model) =====")
549 print(f"R-squared: {r_squared_log:.3f}")
550 print(f"Adj. R-squared: {adj_r_squared_log:.3f}")
551 print(f"F-statistic: {model_log.fvalue:.2f}")
552 print(f"F p-value: {model_log.f_pvalue:.2e}")
553
554 print("\nAll coefficients (log scale):")
555 for name, val in coeffs_log.items():
556     print(f"  {name:35s} {val: .4f}")
557
558 # Calculate approximate percentage effects - SIMPLIFIED VERSION
```

```

559 print("\nApproximate percentage effects (relative to baseline):")
560
561 # Manual mapping based on what we see in the coefficients
562 for name, coef in coeffs_log.items():
563     if "Sector" in name and name != "C(Sector)":
564         pct_effect = (np.exp(coef) - 1) * 100
565         # Extract the actual sector name more carefully
566         if '[, in name and ,]' in name:
567             sector_name = name.split(', [,')[0].split(',')[-1]
568         elif 'T.' in name:
569             sector_name = name.split('T.')[0].strip(',')
570         else:
571             sector_name = name.split('.').strip()[-1] if '.' in name else name
572         print(f"    Sector {sector_name:20s} {pct_effect:+.1f}%")
573
574     elif "Region" in name and name != "C(Region)":
575         pct_effect = (np.exp(coef) - 1) * 100
576         # Extract the actual region name more carefully
577         if '[, in name and ,]' in name:
578             region_name = name.split(', [,')[0].split(',')[-1]
579         elif 'T.' in name:
580             region_name = name.split('T.')[0].strip(',')
581         else:
582             region_name = name.split('.').strip()[-1] if '.' in name else name
583         print(f"    Region {region_name:20s} {pct_effect:+.1f}%")
584
585 if 'TimeNumeric' in coeffs_log:
586     print(f"\nTime trend: {coeffs_log['TimeNumeric']:.6f} (log units per
587 day)")
588     print(f"Approximate annual trend: {((np.exp(coeffs_log['TimeNumeric']) *
589 365) - 1) * 100:.2f}% per year")
590
591 # Model comparison
592 print(f"\n===== MODEL COMPARISON =====")
593 print(f"Original data R : {model_orig.rsquared:.3f}")
594 print(f"Log-transformed R : {model_log.rsquared:.3f}")
595 print(f"Recommend using {'LOG-TRANSFORMED' if model_log.rsquared >
596 model_orig.rsquared else 'ORIGINAL'} model")
597
598 print(f"\nAll figures have been saved to: {FIGURES_PATH}")
599 print("\n===== ANALYSIS COMPLETE =====")
600
601 # =====
602 # EXPORT ALL STATISTICAL RESULTS TO A TXT FILE
603 # =====

```

```
601
602 report_path = FIGURES_PATH / "final_statistics_report.txt"
603
604 with open(report_path, "w", encoding="utf-8") as f:
605
606     f.write("=====\\n")
607     f.write(" FINAL STATISTICAL REPORT\\n")
608     f.write("=====\\n\\n")
609
610     # ----- BASIC SUMMARY STATISTICS -----
611     f.write("==== BASIC SUMMARY STATISTICS =====\\n\\n")
612     f.write("Original Values:\\n")
613     f.write(df_long[["Value"]].describe().to_string())
614     f.write("\nLog-Transformed Values:\\n")
615     f.write(df_long[["Log_Value"]].describe().to_string())
616     f.write("\n\\n")
617
618     # ----- CATEGORY COUNTS -----
619     f.write("==== CATEGORY COUNTS =====\\n\\n")
620     f.write("Sector counts:\\n")
621     f.write(df_long[["Sector"]].value_counts().to_string())
622     f.write("\nRegion counts:\\n")
623     f.write(df_long[["Region"]].value_counts().to_string())
624     f.write("\n\\n")
625
626     # ----- GROUPED SUMMARY -----
627     f.write("==== GROUPED SUMMARY BY REGION & SECTOR =====\\n\\n")
628     f.write(group_summary.to_string(index=False))
629     f.write("\n\\n")
630
631     # ----- REGRESSION OUTPUTS -----
632     f.write("==== MULTIPLE LINEAR REGRESSION (LOG MODEL) =====\\n\\n")
633     f.write(model_log.summary().as_text())
634     f.write("\n\\n")
635
636     f.write("==== ANOVA TABLE (TYPE II) - LOG MODEL =====\\n\\n")
637     f.write(anova_results_log.to_string())
638     f.write("\n\\n")
639
640     f.write("==== MULTIPLE LINEAR REGRESSION (ORIGINAL DATA) =====\\n\\n")
641     f.write(model_orig.summary().as_text())
642     f.write("\n\\n")
643
644     # ----- MODEL INTERPRETATION -----
645     f.write("==== MODEL INTERPRETATION =====\\n\\n")
```

```

646 f.write(f"R-squared (log model): {model_log.rsquared:.4f}\n")
647 f.write(f"Adj. R-squared (log model): {model_log.rsquared_adj:.4f}\n")
648 f.write(f"F-statistic: {model_log.fvalue:.4f}, p-value: {model_log.
f_pvalue:.4e}\n\n")
649
650 f.write("Coefficients (log model):\n")
651 for name, value in coeffs_log.items():
652     f.write(f" {name:35s} {value:.6f}\n")
653 f.write("\n")
654
655 f.write("Approximate percentage effects:\n")
656 for name, coef in coeffs_log.items():
657     if "Sector" in name and name != "C(Sector)":
658         pct = (np.exp(coef) - 1) * 100
659         f.write(f" Sector {name}: {pct:+.2f}%\n")
660     elif "Region" in name and name != "C(Region)":
661         pct = (np.exp(coef) - 1) * 100
662         f.write(f" Region {name}: {pct:+.2f}%\n")
663 f.write("\n")
664
665 if "TimeNumeric" in coeffs_log:
666     annual_trend = (np.exp(coeffs_log["TimeNumeric"] * 365) - 1) * 100
667     f.write(f"Time trend (per day, log-scale): {coeffs_log[",
668 TimeNumeric']:.8f}\n")
669     f.write(f"Approx annual trend: {annual_trend:.4f}%\n\n")
670
671 # ----- MODEL COMPARISON -----
672 f.write("===== MODEL COMPARISON =====\n\n")
673 f.write(f"Original model R : {model_orig.rsquared:.4f}\n")
674 f.write(f"Log-transformed model R : {model_log.rsquared:.4f}\n")
675 better = "LOG-TRANSFORMED" if model_log.rsquared > model_orig.rsquared
else "ORIGINAL"
676 f.write(f"Recommended model: {better}\n\n")
677
678 f.write("===== END OF REPORT =====\n")
679 f.write("===== END OF REPORT =====\n")
680
681 print(f"\nTXT report successfully written to:{report_path}")

```

Listing 1: Data Processing and Analysis Script